

Make your datacenter flex with the grid.

FlexCast predicts Texas grid power for every hour of the coming week, then re-times an AI campus's training around that read. Same compute, smaller bill, lighter grid.

THE PROBLEM

Texas wants the AI buildout and the grid cannot absorb it. A 500 MW campus draws like a mid-size city, flat, day and night, and new loads wait years in the interconnection queue. In 2025 Texas passed SB6: loads over 75 MW must be able to shed for the grid to connect at all. And the power bill is a casino: the same megawatt-hour costs about \$20 on a quiet night and \$3,000 in a scarcity hour.

THE INSIGHT

AI training is the one giant industrial load that can move. Training is deadline-shaped: what matters is that a run finishes by Friday, not which hours it runs. It can pause at a checkpoint, slow its chips, or wait for midnight, while the inference floor serving real users is never touched. SB6 makes that flexibility the price of admission; FlexCast makes it pay every week.

WHAT FLEXCAST DOES

1 · Predict

Ten gradient-boosted models (LightGBM) trained on 5.5 years of real ERCOT prices, demand and weather. Two causally separated horizons, quantile bands for honest uncertainty, calibrated spike probabilities.

3 · Act

The plan lands as a day-by-day playbook: train hard tonight, hold at the floor 3 to 7 PM, battery covers the peak. A live reflex layer answers real-time spikes, pausing only what it can provably repay.

2 · Plan

A linear program re-times every job around the forecast: training into cheap valleys, battery charged cheap and spent into peaks, billing-peak afternoons ducked. Deadlines and the serving floor are hard constraints.

4 · Prove

Every Monday for 20 held-out months, the replay forecast with only what was knowable, committed the plan, and let real prices grade it. No hindsight anywhere in the loop.

SCORED ON 20 MONTHS OF DATA THE MODEL NEVER SAW · REAL PRICES, REAL 4CP PEAKS

\$30.5M

saved vs never flexing

12.2%

off the power bill

510/510

deadlines kept

454 MW

avg draw, frontier-hot

Honest by construction: models trained through 2024 and graded on 2025 to 2026; flexing costs (checkpoint friction, battery wear, throughput loss) are stated estimates; the proof runs the campus realistically hot at 87 percent fleet utilization. Every number on screen traces to a file the pipeline writes.